

Week 6

Hash Tables

purpose: maintain a set of stuff with
all operations in $O(1)$ time!

Operations:

- insert
- delete
- lookup

} all in $O(1)$ time! *

* when properly implemented
and data is not pathological

"dictionary"

but a hash table doesn't support
the order!

Applications:

△ De-duplications

given: a "stream" of objects

goal: ignore duplicates (keep track of
unique objects)

solution: when a new object x arrives

- look x up in hash table H
- if not found, insert x into H .

Application:

△ The 2-SUM problem

input: unsorted array A of n integers
target sum t

goal: determine whether or not there are
two numbers x, y in A with

$$x + y = t$$

naïve solution: $O(n^2)$ time

better: 1) sort A ($O(n \log n)$ time)

2) for each x in A look for
 $t - x$ in A via binary search
($O(n \log n)$ time)

more

better: 1. insert elements of A into
hash table H ($O(n)$ time)
2. for each x in A ,
lookup $t - x$ ($O(n)$ time)

Also

△ historical application: symbol tables in
compilers

△ Blocking network traffic (blacklist)

△ Search algorithms (game tree exploration, etc)

△ etc.

high-level idea

Setup: universe U (all IP addresses, all names, etc)
generally, really, really big

Goal: want to maintain set $S \subseteq U$

Solution: pick n - number of buckets
(assume $|S|$ doesn't change much)

1. choose a hash function
 $h: U \mapsto \{0, 1, \dots, n-1\}$
2. use array A of length n , store x
in $A[h(x)]$

Birthday paradox

Consider n people with random birthdays.
How large does n need to be before there
is at least 50% chance that two people
have the same birthday?

23 (50%) 57 (99%)

↳ Collisions!

Collision: distinct $x, y \in U$
such that $h(x) = h(y)$

Solution #1

(separate) chaining

- keep linked list in each bucket
- given a key/object x , perform Insert/Delete/
lookup in the list in $A[h(x)]$

↙
returns a list

Solution #2

open addressing (only one object per bucket)

hash function now specifies probe
sequence $h_1(x), h_2(x), \dots$

- keep trying until we find an open slot

examples: linear probing
double hashing

- A Good hash function

note: in hash table with chaining,
 $O(\text{list len})$ for insert/delete

all objects are in the same bucket ————— could be anywhere from (m/n) to m for m objects equal-len list

point: performance depends on the choice of hash function!

So, "good" hash function should:

- spread data out \Rightarrow lead to good performance
- easy to store / be fast to evaluate

example:

keys - phone numbers (10-digits)

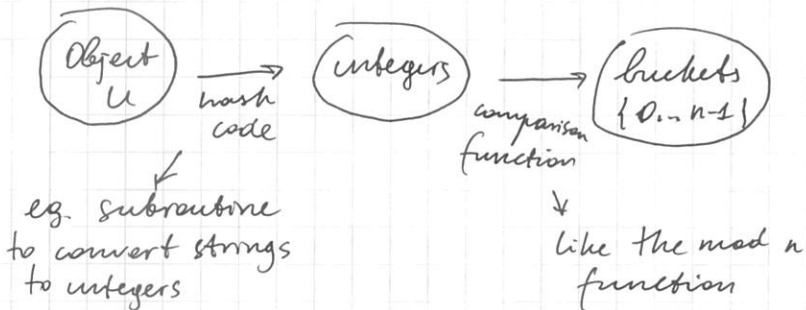
- terrible hash: $h(x)$ - 1st 3 digits of x
- mediocre hash: $h(x)$ - last 3 digits of x

example:

keys - memory locations

- bad hash: $h(x) = x \bmod 1000$
(all odd buckets guaranteed to be empty)

Quick-and-Dirty hash function



How to choose $n = \#$ of buckets

1. Choose n to be prime (without constant factor of $\#$ of objects in table)
2. Not too close to a power of 2 or 10

The load of a hash table

the load factor

$$\alpha = \frac{\text{number of objects}}{\text{number of buckets}}$$

- for good hash table performance, need to control load
- and a good hash function

Pathological data sets

super- clever user-defined hash function does not guarantee to spread evenly

For every hash function there exists a pathological dataset

Solutions :

1. Use a cryptographic hash function (SHA-2, etc)

2. Use randomization

design a family H of hash functions such that for all data sets S "almost all" functions $h \in H$ "spread S out" "pretty evenly"

Universal Hash Functions

~~Overview~~ Definition

- let H be a set of hash functions from U to $\{0, 1, \dots, n-1\}$

H is universal if and only if:
for all $x, y \in U$ ($x \neq y$)

$$\Pr_{h \in H} [x, y \text{ collide}] \leq \frac{1}{n}$$

i.e. $h(x) = h(y)$

when h is chosen uniformly at random from H .

↑
(number of buckets)

example:

hashing IP addresses

let U = IP addresses of the form
 (x_1, x_2, x_3, x_4) with each $x_i \in \{0, 1, \dots, 255\}$

let n = a prime

construction: define one hash function
 h_a per 4-point

$$a = (a_1, a_2, a_3, a_4)$$

with each $a_i \in \{0, 1, \dots, n-1\}$

h_a = IP address \rightarrow buckets by

$$h_a(x_1, x_2, x_3, x_4) = \begin{pmatrix} a_1 x_1 + a_2 x_2 + \\ a_3 x_3 + a_4 x_4 \end{pmatrix} \bmod n$$

n^4 such functions

\Downarrow

$$H = \{h_a \mid a_1, a_2, a_3, a_4 \in \{0, \dots, n-1\}\}$$

$$h_a = \sum a_i x_i \bmod n$$

this family is universal

Bloom Filters

reason: fast inserts and lookups

comparison to hash tables:

pros

- more space efficient

cons

- can't store an associated object
- no deletions
- small false positive probability
(i.e. may say x has been inserted, but it wasn't)

Applications

- original - early spellcheckers
- canonical: list of forbidden passwords
- modern: network routers
(limited memory, need to be super fast)

Under the hood:

Ingredients:

1. array of n bits ($\frac{n}{|S|}$ - number of bits per object in data set S)
2. k -hash functions $h_1 \sim h_k$ (k - small constant)

Insert(x):

for $i = 1 \dots k$

set $A[h_i(x)] = 1$

Lookup(x)

True if $A[h_i(x)] == 1$ for every
 $i = 1 \dots k$

note: no false ^{negatives} ~~positives~~ (if x was
inserted, Lookup(x) guaranteed to
succeed)

but: false positives if all k $h_i(x)$'s
already set to 1 by other
insertions