

## Week 6

### Predictions

- motivation:
- ~~key~~ ~~steps~~ steps in predictive studies
- error measure

kaggle.com - prediction contests

### Steps

1. Find the right data
2. Define your error rate
3. Split data into
  - training
  - testing
  - validation (optional)
4. On the training set pick features
5. On the training set pick prediction function
6. On the training set - cross-validate

if no validation:

apply a 1x to test set

if validation

apply to test set and refine

apply 1x to validation

True/False Positives (for binary classification)

Positive = identified

Negative - rejected

	Positive	Negative
True	correctly identified	<del>correctly</del> rejected
False	incorrectly identified	incorrectly rejected

e.g.

	Positive	Negative
True	Sick people correctly diagnosed	Healthy identified as healthy
False	Healthy people diagnosed as sick <del>healthy</del>	sick people identified as healthy

Error rate:

Positive predictive value =

$$\frac{\sum TP}{\sum \text{test outcome positive}} \leftarrow \begin{array}{l} \text{true value} \\ \text{test value} \end{array}$$

Negative predictive value

$$\frac{\sum TN}{\sum \text{test outcome negative}} \leftarrow \begin{array}{l} \text{true} \\ \text{test} \end{array}$$

$$\text{Sensitivity} = \frac{\sum TP}{\sum \text{Condition positive}} \leftarrow \begin{array}{l} \text{true value} \\ \text{test value} \end{array}$$

( $\frac{\text{sick}}{\text{test sick}}$ )

$$\text{Specificity} = \frac{\sum TN}{\sum \text{Condition Negative}} \leftarrow \left( \frac{\text{Healthy}}{\text{test Healthy}} \right)$$

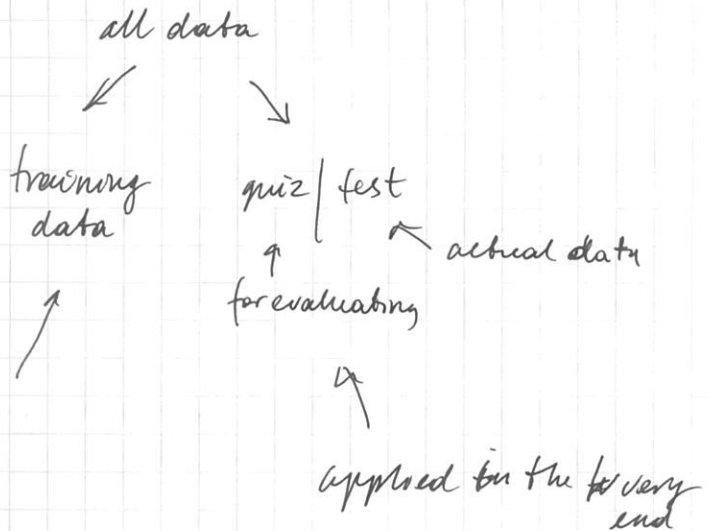
False positive - Type 1 error

False negative - Type 2 error

## Common error measures

- mean square error  
continuous data, sensitive to outliers
- median absolute deviation  
continuous data, often more robust
- sensitivity (recall)  
if you want few ~~miss~~<sup>ed</sup> positives
- specificity  
if you want few negatives called positives
- accuracy  
weight false positives/negatives equally

# Study design



## key issues:

- accuracy
- over fitting (you fit well to training set, but don't to other samples)
- interpretability
- computational speed

## Resources:

- practical machine learning
- elements of statistical learning
- coursera machine learning
- machine learning for hackers

## Cross-validation

how to estimate error rate for your predictive functions

key ideas:

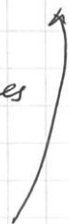
- sub-sampling the training data
- avoiding overfitting
- making predictions generalizable

Goal of cross-validation - how well your predictions will work on different samples

- accuracy on training set is optimistic
- a better estimate comes from an independent set

but we cannot ~~test~~ use the test set when building the model - or it becomes the part of the training set

test set  
accuracy



## basic approach




1. Use the training set
2. Split it into training / test sets
3. Build a model on the training set
4. Evaluate on the test set
5. Repeat and average the estimate errors

Used for:

- picking variables to include in a model
- picking the type of prediction function to use
- picking the parameters in the prediction function
- comparing different predictors

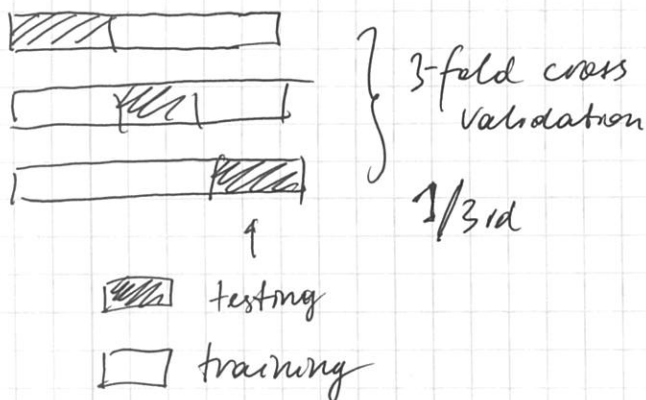
Ways:

- random subsampling

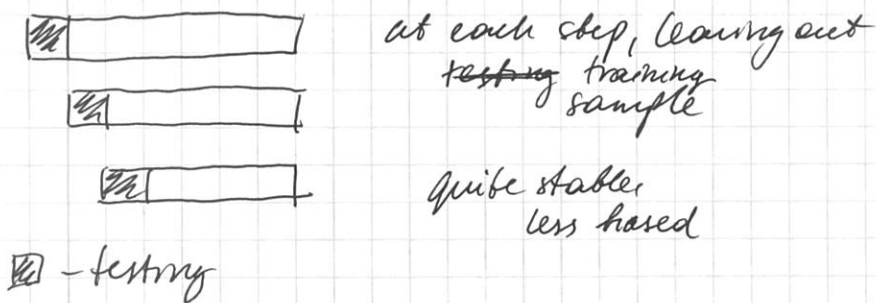


← do over and over again

## K-fold



## leave one out



## Notes

- training sets and testing sets must come from the same population
- sampling should be designed to mimic real patterns

(sample chunks of times - not just random sampling for time series)



# Predicting with regressions

- lm or glm -

predict new values with coefficients

pros:

easy to implement  
easy to interpret

cons:

poor performance in non-linear settings

predicting:

$$\text{coef}(\text{lm1})[1] + \text{coef}(\text{lm2})[2] * 80$$

↑

what you want  
to predict

↙ equivalent

predict function

Calculating training set / test set errors

# RMSE

$$\sqrt{\text{sum}(\text{lm1}\$fitted - \overset{\text{test}}{\text{train}\$erruptrons})^2)}$$

↑  
var we're  
predicting

$\sqrt{\text{sum}(($

(how but not why)

$\text{predict}(\text{lm1}, \text{newdata} = \text{test} - \text{test\$eruptions})^2$

$)$

↓  
error's slightly larger

prediction intervals

95%-prediction

good for normally distributed data

Choosing a cut-off. (re-substitution)

a value at which you get the smallest amount of errors

Comparing models

$\text{cost} = \text{function}(\text{urn}, \text{pred} = 0) \text{ mean}(\text{abs}(\text{urn} - \text{pred}) > 0.5)$

when the error is high enough

library(boot)

error est. function

cv1 = cv.glm(ravens data, glm1, cost, K=3)

↑  
regression  
model

↑  
number of  
folds

cv1\$delta

↑  
the less the  
better

↑  
for checking  
what regression model is  
better

## Predicting with Trees

Better at capturing non-linearity

key ideas:

- iteratively split variables into groups
- split where maximally predictive
- evaluate "homogeneity" within each branch  
(how similar the objects are  
within the groups)  
and maximize it in each group
- fitting multiple trees often works better

Pros:

- easy to implement
- easy to interpret
- better performance in non-linear settings

Cons:

- without pruning/cross-validation can lead to overfitting
- harder to estimate uncertainty
- results may be variable

example - decision tree

Basic algorithm:

1. Start with all variables in <sup>one</sup> group
2. Find the variable/split that best separates the outcomes (in homogeneous)
3. Divide the data into two groups ("leaves") on that split ("node")
4. Within each split, find the best variable/split that separates the outcomes
5. Continue until the groups are small enough or sufficiently "pure"

~~predict~~

## Example

data (iris) — need to classify it

library (tree) outcome

tree1 = tree (Species ~ Sepal.Width + Petal.Width)

summary (tree1)

↓

says the number of terminal nodes: 5


Residual <sup>(measure of impurity)</sup> mean deviance:  $0.204 = 29.6 / 145$

Misclassification error rate:  $0.0333 = 5 / 150$   
(how often you misclassify)

plot (tree1)

text (tree1)

partition.tree (tree1)



pred1 = predict (tree1, newdata) — returns probabilities

pred1 = predict (tree1, newdata, type = "class")

partition.tree (tree1, "Species", add=T) — returns not probabilities, but a class

## Pruning

if we built a tree and want to cross-validate it


cv. tree - for this

```
plot(cv.tree(treeVars, FUN = prune.tree, method =  
      "misclass"))
```

↓  
number of misclassification  
errors

cv.tree(treeVars)

default method is deviance



as the model increases  
in size

- start with 35 misclassification
- between 6-9: fewer misclassification
- then we have misclassification again

↑  
What happens to different sizes of  
the model

`prune.tree (tree1, best = 4)`

give me the best tree that has exactly  
4 terminal leaves / nodes

And it gives a smaller subset of the tree -

And this smaller set will get a smaller error  
rate

(`prune - odfprune,`  
`no pruning [general]`)

Resubstitution Error

`table (orig, predict(pruneTree, type = "class"))`

↓

shows a table

works better than not pruned trees.