

Week 7

Smoothing

for observing non-linear trends - and for being able to predict those trends in other observations

it tries to fit a curvy smooth line -

but there is a risk of overfitting and it's hard to interpret.

in R -

filter - smooths data

as.vector = (filter(cd4\$time, filter = rep(1, 200))) / 200

averaging - is smoothing (or weighted sum)

↑
"moving average"

sum of weights
should be 1

↓
can use  weight

the closer to
the middle, the
more weight

lowess (loess) \leftarrow locally-weighted
scatter-plot smoothing

lw1 = loess(cd4 ~ time, data = c4)

lines(time, lw1\$fitted, ...)

span - number of points used when smoothing

span = 0.1 - 0.1 of data is used for
calculating

0.25 - a quarter of the data is used

as the span increases, the line becomes
more smooth

by default, loess calculates the span
automatically

predicting:

pred1 = predict(lw1, new.data = ..., se = T)

\uparrow
calculates the
standard error

Splines

$$Y_i = b_0 + \sum_{k=1}^K b_k s_k(x_i) + \epsilon_i$$

Y_i - outcome of the observation

b_0 - intercept term

b_k - coefficient for k^{th} spline function

x_i - covariate for i^{th} observation

ϵ_i - error

library(splines)

ns1 = ns(cd4\$ time, df = 3)

↑
natural
cubic splines

↑
degrees of freedom
(number of functions to
be applied to the variable)

lm1 = lm(cd4 ~ ns1)

summary(lm1)

↑
spline matrix

Notes:

cross-validation is important

The bootstrap

for estimating standard errors
for improving predictions

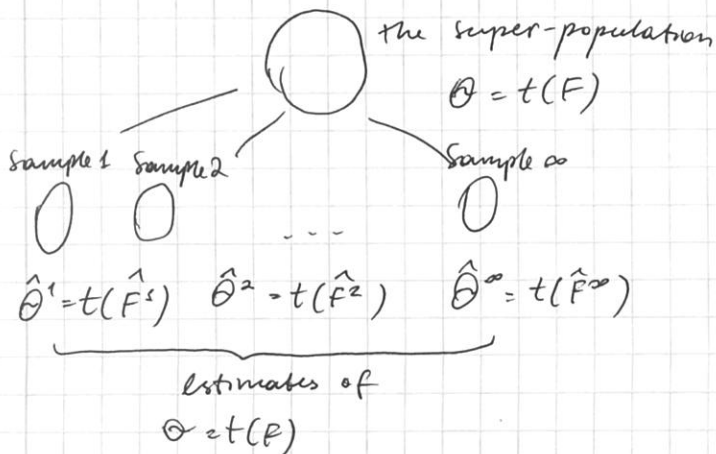
key idea:

- treat the sample as population -
and then perform analysis as if
you taken a sample from population

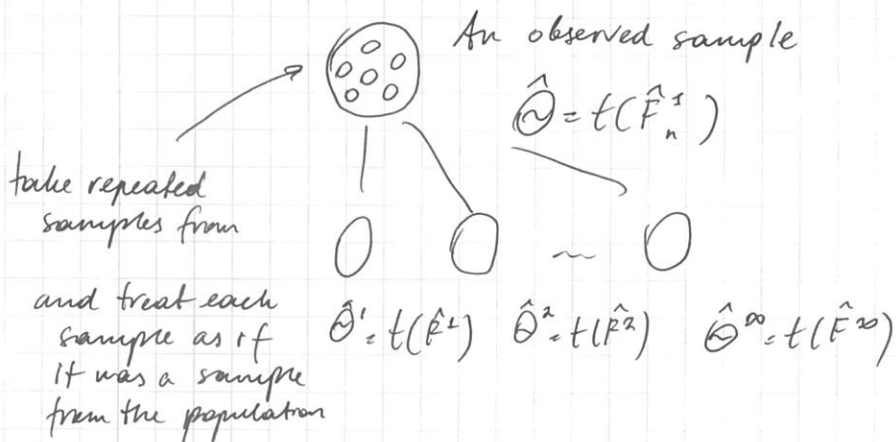
good for

- calculating standard errors
- forming confidence intervals
- performing hypothesis testing
- improving predictors

The central dogma of statistics



bootstrap



this idea is quite powerful and works very well

bootMean = boot(x, meanFunc, 1000)

original data

number of iterations

meanFunc = function(x, i) {
 mean(x[i])
}

dataset set of indexes

bootMean \$t

for calculating confidence intervals

library(boot) data(nuclear)

nuke.lm = lm(log(cost) ~ date, data = nuclear)

without making assumptions about normal distribution of the data set or any other assumptions

results = boot(data = nuclear, statistics = bs,
 R = 1000, formula =
 ~log(cost) ~ date)

number of iterations

bs = function(data, indices, formula,

d = data[indices,]

fit = lm(formula, data = d)

return(coef(fit))

}

Combination
of boot and bs function
generates 1000 linear fits
where we resample the data set
with replacements

boot.ci(results)

↑
bootstrap confidence intervals

Bootstrapping from a model

resrd = rstudent(nuke.lm) -- calculate
Student's residuals

fit = fitted(lm(log(cost) ~ 1;
data = nuclear)) ← fitted value if
we only included
the intercept term
(when you fit a straight
line to the data)

newNuc = cbind(nuclear, resrd = resrd,
fit0 = fit0)

bs = function(data, indices) {
 coef(glm(data\$fit0 + ^{outcome} data\$resrd[indices] ~
 data\$date, data = data))
}

results = boot(data = newNuc, statistic = bs,
R = 1000)

Things you can't bootstrap

- map

notes

- useful for complicated statistics
- careful near boundaries
- careful with non-linear functions

resources

~~An introduction to the bootstrap,
from the stude r~~

Bootstrapping for predictions

newdata = data.frame(data = seq(65, 72, length = 100))
nuclear = cbind(nuclear, resid = rstudent(nuke.lm),
fit = fitted(nuke.lm))

↑
fitted value
from nuke.lm

↑
residual value
from nuke.lm

dates between 65 and 72

nuke.boot = boot(nuclear, nuke.fun, R = 1000,
newdata = newdata)

nuke.fun = function (data, inds, ^{randomised bootstrap indices} newdata) {

lm.b = lm (fit + resid [inds] ~ date, data = data)

↑
bootstrap-sampled
noise

it tries to generate the
data that has the
same trend observed in
the original dataset

pred.b = predict (lm.b, newdata)

↑
predicting values for new data

return (pred.b)

}

nuke.boot \$t - 100 replications for
each of values from 65 to 72
(in columns)

```
pred = predict(mike.lm, newdata)
```

↑
our predicted value for new data set


```
predStds = apply(mike.boot$t, 2, sd)
```

↑
standard deviations for each of
those 100 values

```
plot(newdata$date, col = "black", ylim = c(0, 10))
```

```
lines(newdata$date, pred + 1.96 * predStds, col = "red")
```

```
lines(newdata$date, pred - 1.96 * predStds, col = "red")
```



Bootstrap aggregating

for improving prediction accuracy

basic ideas

1. Resample cases and recalculate predictions
2. average or majority vote

(have a look at
wiki)

it reduces variances
and more useful for non-linear functions

```
library(ElemStatLearn)  
data(ozone, package = "ElemStatLearn")
```

```
ozone = ozone[order(ozone$ozone),]
```

```
U = matrix(NA, nrow = 10, ncol = 155)
```

```
for (i in 1:10) {
```

```
  ss = sample(1:dim(ozone)[1], replace = T)
```

```
  ozone0 = ozone[ss,] ← subsample
```

```
  loess0 = loess(temperature ~ ozone,  
                 data = ozone0, span = 0.2)
```

```
  U[i,] = predict(loess0, newdata =  
                  data.frame(ozone = 1:155))
```

```
}
```

Bagged
loess

```
lines(1:155, apply(U, 2, mean), col = "red")
```

↑
bagged less line

Bagged trees

basic idea

- resample data
 - recalculate tree
 - average/mode of predictors
-
- more stable
 - may not be as good as random forests

library (imed)

bag tree = bag song (species ~, wab = T)

↑
some values
are left out

returns a bunch of trees
with different classifications

Random forests

1. Bootstrap samples
2. At each step, bootstrap variables
3. Grow multiple trees and vote

pros

- accuracy

cons

- speed
- interpretability
- overfitting

library(randomForest)

forestIris = randomForest (Species ~ Petal.Width +
Petal.Length, data = iris, prox = T)

getTree (forestIris, k = 2) - returns a tree

```
iris.p = classCenter (iris[, c(3, 4)],  
                      iris$Species, forestIris$prox)
```

↓

draws "centers of clusters"

combine (...)

↑

combines random forests into one

prediction: the same

predict (forest Brs, newdata)

Notes:

- bootstrapping is useful for non-linear models
- care should be taken to avoid overfitting
- out of bag estimates are efficient estimates of test error

Combining Predictors

key ideas

- you can combine classifiers by averaging/voting



improves accuracy
but reduces interpretability

Base intuition

Suppose we have 5 completely independent classifiers

if accuracy is 70% for each

$$10 \times (0.7)^3 (0.3)^2 + 5 \times (0.7)^4 (0.3)^1 + (0.7)^5 =$$

10 ways when 3 classifiers are right and 2 wrong \uparrow 5 ways when 4 right and 1 wrong

= 83.7% accuracy

which is higher than accuracy of any of the individuals

with 101 independent \Rightarrow 99.9%

Approaches for combining:

- bagging
- boosting
- combining (in weighted/unweighted fashion) — avg or majority different classifiers

$$\text{Combine 1} = \text{predict}(\text{lm1}, \text{data} = \text{test Data}) / 2 + \\ + \text{predict}(\text{tree1}, \text{data} = \text{test Data}) / 2$$